

CS634 Final Term Project Report

Yash Shah
CS634 – 101
Dr. Yasser Abdullaah

Telecom Churn Prediction Data Using Random Forest, Conv 1D, and SVM
Nov 10, 2025

1. Introduction

The purpose of this project is to build, evaluate, and compare multiple machine learning models for binary classification. Specifically, the goal is to predict customer churn (whether a telecom customer will stay or leave), based on demographic, account, and service-related features. This helps businesses identify at risk customers and take proactive retention actions.

Binary classification is a supervised learning task where the model predicts one of two possible outcomes (e.g., Yes/No, Churn/Not Churn, Fraud/Not Fraud). It's important because many real-world decisions rely on such two-class predictions which helps enable targeted marketing, risk management, medical diagnosis, and customer retention strategies. Accurate binary classification helps optimize business decisions and resource allocation.

Dataset: Telco Customer Churn Dataset which contains customer demographics, service subscriptions, billing details, and churn status (Yes/No).

Algorithms Implemented:

Random Forest (mandatory): An ensemble learning method that builds multiple decision trees and averages their results for robust predictions.

Conv1D (Deep Learning Model): A 1D Convolutional Neural Network that captures local feature patterns and nonlinear relationships in tabular data.

Support Vector Machine (SVM): A classic machine learning algorithm that finds the optimal boundary between churners and non-churners using kernel functions.

2. Dataset

Name: Telecommunication Churn Dataset - Kaggle

Link: <https://www.kaggle.com/datasets/blastchar/telco-customer-churn>

The Telco Customer Churn dataset contains information about customers from a telecommunications company, including demographic details, account information, and subscribed services.

Number of rows: 7,043

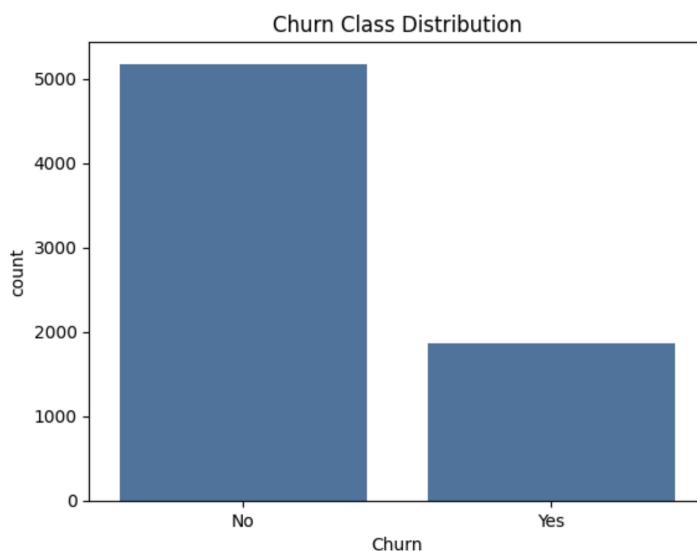
Number of features: 20 independent variables

Target variable: Churn - a binary variable indicating whether a customer has left the company (Yes = 1, No = 0).

Key attributes include customer tenure, contract type, payment method, monthly and total charges, and the presence of services such as internet, phone, and technical support.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   customerID  7043 non-null   object  
 1   gender       7043 non-null   object  
 2   SeniorCitizen 7043 non-null   int64  
 3   Partner      7043 non-null   object  
 4   Dependents   7043 non-null   object  
 5   tenure       7043 non-null   int64  
 6   PhoneService 7043 non-null   object  
 7   MultipleLines 7043 non-null   object  
 8   InternetService 7043 non-null   object  
 9   OnlineSecurity 7043 non-null   object  
 10  OnlineBackup   7043 non-null   object  
 11  DeviceProtection 7043 non-null   object  
 12  TechSupport    7043 non-null   object  
 13  StreamingTV    7043 non-null   object  
 14  StreamingMovies 7043 non-null   object  
 15  Contract       7043 non-null   object  
 16  PaperlessBilling 7043 non-null   object  
 17  PaymentMethod   7043 non-null   object  
 18  MonthlyCharges 7043 non-null   float64 
 19  TotalCharges   7043 non-null   object  
 20  Churn          7043 non-null   object  
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
None
```

Unique values per column:		
customerID	7043	
gender	2	
SeniorCitizen	2	
Partner	2	
Dependents	2	
tenure	73	
PhoneService	2	
MultipleLines	3	
InternetService	3	
OnlineSecurity	3	
OnlineBackup	3	
DeviceProtection	3	
TechSupport	3	
StreamingTV	3	
StreamingMovies	3	
Contract	3	
PaperlessBilling	2	
PaymentMethod	4	
MonthlyCharges	1585	
TotalCharges	6531	
Churn	2	
dtype:	int64	



Preprocessing Steps

Handling Missing Values: The TotalCharges column contained blank strings, which were converted to numeric values using pd.to_numeric(errors='coerce') and filled with the median. All other columns had no missing values.

Encoding Categorical Variables: Categorical features (e.g., Gender, Contract, PaymentMethod) were transformed using One-Hot Encoding to convert text labels into numerical format suitable for model training.

Normalization / Standardization: Numerical features such as tenure, MonthlyCharges, and TotalCharges were standardized using StandardScaler to ensure all variables were on comparable scales, improving model convergence.

Class Balance Handling: The dataset was slightly imbalanced (~26.5% churners vs. 73.5% non-churners). To address this, models were trained using class weights = 'balanced' (for Random Forest and SVM) and weighted loss (for Conv1D), ensuring fair learning from both classes.

3. Algorithms Overview

1. Random Forest

Rationale:

Random Forest was chosen as the baseline model because it is robust, handles both numerical and categorical features well, and performs automatic feature selection. It is also resistant to overfitting and provides strong predictive performance even with minimal parameter tuning.

How it works:

Random Forest is an ensemble learning algorithm that builds multiple independent decision trees on random subsets of data and features. Each tree votes for a class (churn or not), and the final prediction is made by majority vote. Averaging across many trees reduces variance and improves generalization.

2. Deep Learning Model – Conv1D

Rationale:

A 1D Convolutional Neural Network (Conv1D) was selected as the deep learning model because it can effectively capture local feature patterns and interactions within tabular data. Unlike dense networks, Conv1D applies convolution filters across the feature space, detecting meaningful combinations (e.g., between tenure, billing, and service types) that influence churn behavior.

Why it suits this task:

Handles complex nonlinear relationships among features.

Learns high-level representations automatically, reducing reliance on manual feature engineering.

Works efficiently on relatively small datasets compared to deeper architectures like LSTM.

3. Classic Machine Learning Model – Support Vector Machine (SVM)

Rationale:

SVM with an RBF (Radial Basis Function) kernel was selected as the traditional ML algorithm for its ability to model nonlinear decision boundaries effectively.

Strengths:

Performs well in high-dimensional feature spaces.
Maximizes the margin between classes, improving generalization.
Works effectively with smaller datasets.

Limitations:

Computationally expensive on large datasets.
Requires careful parameter tuning (C and gamma).

4. Implementation

Programming language: Python.

Development environment: Jupyter Notebook and .py scripts.

In VsCode

1. In your project folder, open a terminal and run:

```
python -m venv venv
```

```
venv\Scripts\activate # for Windows
```

```
pip install numpy pandas scikit-learn tensorflow keras matplotlib seaborn
```

OR

```
pip install requirements.txt (from my Github link)
```

2. To run standalone code:

```
python churn_models.py
```

3. To run Jupyter notebook:

```
jupyter notebook .\Yash_Shah_Final.ipynb
```

Dataset Loading Screenshot:

```
df = pd.read_csv(CSV_PATH)

print(df.head())

***   customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0    7590-VHVEG  Female           0      Yes        No         1        No
1    5575-GNVDE   Male            0       No        No        34       Yes
2    3668-QPYBK   Male            0       No        No         2       Yes
3    7795-CFOCW   Male            0       No        No        45        No
4    9237-HQITU  Female           0       No        No         2       Yes

    MultipleLines InternetService OnlineSecurity ... DeviceProtection  \
0  No phone service             DSL          No  ...          No
1           No                  DSL          Yes  ...          Yes
2           No                  DSL          Yes  ...          No
3  No phone service             DSL          Yes  ...          Yes
4           No  Fiber optic          No  ...          No

    TechSupport StreamingTV StreamingMovies  ... Contract PaperlessBilling  \
0        No          No           No  Month-to-month      Yes
1        No          No           No        One year        No
2        No          No           No  Month-to-month      Yes
3       Yes          No           No        One year        No
4        No          No           No  Month-to-month      Yes

    PaymentMethod MonthlyCharges  TotalCharges  Churn
0  Electronic check        29.85        29.85     No
1      Mailed check        56.95      1889.5     No
2      Mailed check        53.85      108.15    Yes
3  Bank transfer (automatic)  42.30      1840.75     No
4  Electronic check        70.70      151.65    Yes

[5 rows x 21 columns]
```

Training of algorithms:

```
[GLOBAL] RF best params: {'clf__max_depth': 20, 'clf__max_features': 'log2', 'clf__min_samples_leaf': 9, 'clf__min_samples_split': 8, 'clf__n_estimators': 400}
[GLOBAL] SVM best params: {'clf__C': np.float64(21.368329072358772), 'clf__gamma': np.float64(0.0007068974950624604)}
Epoch 1/36
90/90 ━━━━━━━━━━ 3s 12ms/step - loss: 0.6939 - val_loss: 0.6901 - learning_rate: 0.0010
Epoch 2/36
90/90 ━━━━━━━━ 1s 12ms/step - loss: 0.6778 - val_loss: 0.6883 - learning_rate: 0.0010
Epoch 3/36
90/90 ━━━━━━ 1s 13ms/step - loss: 0.6682 - val_loss: 0.6780 - learning_rate: 0.0010
Epoch 4/36
90/90 ━━━━ 1s 13ms/step - loss: 0.6370 - val_loss: 0.5724 - learning_rate: 0.0010
Epoch 5/36
90/90 ━━━━ 1s 8ms/step - loss: 0.6066 - val_loss: 0.6272 - learning_rate: 0.0010
Epoch 6/36
90/90 ━━━━ 1s 7ms/step - loss: 0.5813 - val_loss: 0.5547 - learning_rate: 0.0010
Epoch 7/36
90/90 ━━━━ 1s 7ms/step - loss: 0.5659 - val_loss: 0.6022 - learning_rate: 0.0010
Epoch 8/36
90/90 ━━━━ 1s 8ms/step - loss: 0.5645 - val_loss: 0.5101 - learning_rate: 0.0010
Epoch 9/36
90/90 ━━━━ 1s 8ms/step - loss: 0.5450 - val_loss: 0.5203 - learning_rate: 0.0010
Epoch 10/36
90/90 ━━━━ 1s 8ms/step - loss: 0.5539 - val_loss: 0.5390 - learning_rate: 0.0010
Epoch 11/36
90/90 ━━━━ 1s 8ms/step - loss: 0.5496 - val_loss: 0.5078 - learning_rate: 5.0000e-04
Epoch 12/36
90/90 ━━━━ 1s 8ms/step - loss: 0.5520 - val_loss: 0.5273 - learning_rate: 5.0000e-04
Epoch 13/36
90/90 ━━━━ 1s 7ms/step - loss: 0.5311 - val_loss: 0.5388 - learning_rate: 5.0000e-04
Epoch 14/36
90/90 ━━━━ 1s 8ms/step - loss: 0.5200 - val_loss: 0.5483 - learning_rate: 2.5000e-04
Epoch 15/36
```

Per fold Evaluation Output:

Is included in the Results section.

5. Evaluation Step

A 10-fold cross-validation approach was used to ensure robust model evaluation. The dataset was split into 10 equal parts in each iteration, 9 folds were used for training and 1 fold for testing, rotating until every fold served as the test set once. This reduces variance in performance estimates and provides a more reliable comparison across models.

Metrics Reported (per fold and overall averages):

Confusion-matrix metrics:

TP (True Positives), TN (True Negatives), FP (False Positives), FN (False Negatives)

Performance metrics:

Accuracy, Precision, Recall, F1-score

Rate-based metrics:

FPR (False Positive Rate), FNR (False Negative Rate), Specificity (TNR), Balanced Accuracy

Skill and reliability metrics:

TSS (True Skill Statistic), HSS (Heidke Skill Score)

Probability-based metrics:

ROC curve and AUC (Area Under the Curve)

BS (Brier Score) and BSS (Brier Skill Score)

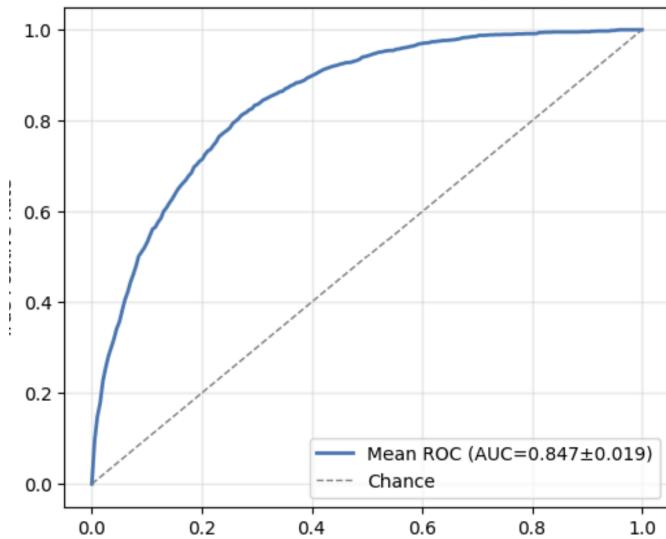
6. Results

Per fold (Random Forest) and ROC curve

==== Random Forest (per fold + average) ===																					
	Fold	TP	TN	FP	FN	P	N	TPR	TNR	FPR	...	Precision	Recall	F1	ErrorRate	TSS	HSS	BS	BSS	BSS_alt	AUC
0	1	146.0	402.0	116.0	41.0	187.0	518.0	0.7807	0.7761	0.2239	...	0.5573	0.7807	0.6503	0.2227	0.5568	0.4936	0.1515	0.2227	0.7773	0.8589
1	2	144.0	399.0	119.0	43.0	187.0	518.0	0.7701	0.7703	0.2297	...	0.5475	0.7701	0.6400	0.2298	0.5403	0.4782	0.1462	0.2499	0.7501	0.8651
2	3	147.0	388.0	130.0	40.0	187.0	518.0	0.7861	0.7490	0.2510	...	0.5307	0.7861	0.6336	0.2411	0.5351	0.4638	0.1578	0.1903	0.8097	0.8533
3	4	141.0	392.0	126.0	45.0	186.0	518.0	0.7581	0.7568	0.2432	...	0.5281	0.7581	0.6225	0.2429	0.5148	0.4518	0.1595	0.1797	0.8203	0.8362
4	5	148.0	406.0	111.0	39.0	187.0	517.0	0.7914	0.7853	0.2147	...	0.5714	0.7914	0.6637	0.2131	0.5767	0.5136	0.1458	0.2528	0.7472	0.8635
5	6	146.0	396.0	121.0	41.0	187.0	517.0	0.7807	0.7660	0.2340	...	0.5468	0.7807	0.6432	0.2301	0.5467	0.4810	0.1522	0.2199	0.7801	0.8574
6	7	138.0	377.0	140.0	49.0	187.0	517.0	0.7380	0.7292	0.2708	...	0.4964	0.7380	0.5935	0.2685	0.4672	0.4044	0.1809	0.0724	0.9276	0.7991
7	8	137.0	423.0	94.0	50.0	187.0	517.0	0.7326	0.8182	0.1818	...	0.5931	0.7326	0.6555	0.2045	0.5508	0.5123	0.1427	0.2687	0.7313	0.8562
8	9	147.0	381.0	136.0	40.0	187.0	517.0	0.7861	0.7369	0.2631	...	0.5194	0.7861	0.6255	0.2500	0.5230	0.4494	0.1603	0.1782	0.8218	0.8385
9	10	132.0	403.0	114.0	55.0	187.0	517.0	0.7059	0.7795	0.2205	...	0.5366	0.7059	0.6097	0.2401	0.4854	0.4410	0.1546	0.2073	0.7927	0.8381
10	Average	142.6	396.7	120.7	44.3	186.9	517.4	0.7630	0.7667	0.2333	...	0.5427	0.7630	0.6338	0.2343	0.5297	0.4689	0.1551	0.2042	0.7958	0.8466

11 rows x 23 columns

Random Forest — Mean ROC (10-fold)

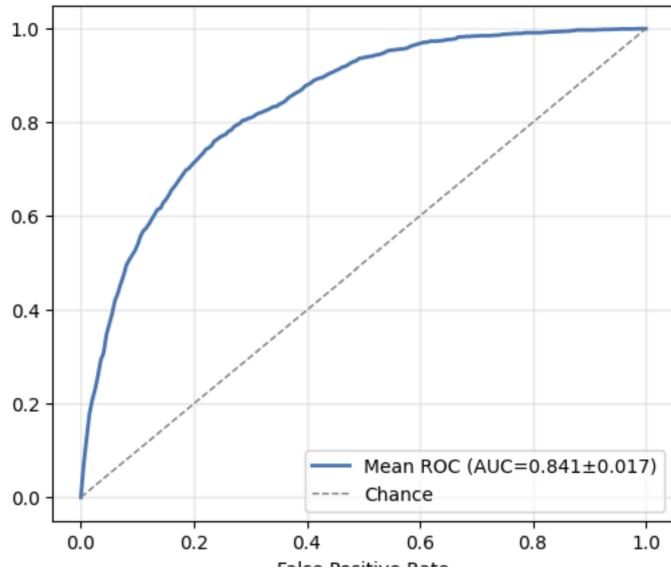


Per fold (SVM) and ROC curve

		SVM (per fold + average)																				
	Fold	TP	TN	FP	FN	P	N	TPR	TNR	FPR	...	Precision	Recall	F1	ErrorRate	TSS	HSS	BS	BSS	BSS_alt	AUC	
0	1	119.0	449.0	69.0	68.0	187.0	518.0	0.6364	0.8668	0.1332	...	0.6330	0.6364	0.6347	0.1943	0.5032	0.5023	0.1389	0.2872	0.7128	0.8566	
1	2	79.0	490.0	28.0	108.0	187.0	518.0	0.4225	0.9459	0.0541	...	0.7383	0.4225	0.5374	0.1929	0.3684	0.4267	0.1559	0.1999	0.8001	0.8465	
2	3	123.0	447.0	71.0	64.0	187.0	518.0	0.6578	0.8629	0.1371	...	0.6340	0.6578	0.6457	0.1915	0.5207	0.5145	0.1470	0.2457	0.7543	0.8559	
3	4	107.0	460.0	58.0	79.0	186.0	518.0	0.5753	0.8880	0.1120	...	0.6485	0.5753	0.6097	0.1946	0.4633	0.4807	0.1583	0.1857	0.8143	0.8291	
4	5	0.0	517.0	0.0	187.0	187.0	517.0	0.0000	1.0000	0.0000	...	0.0000	0.0000	0.0000	0.2656	0.0000	0.0000	0.1527	0.2172	0.7828	0.8649	
5	6	67.0	500.0	17.0	120.0	187.0	517.0	0.3583	0.9671	0.0329	...	0.7976	0.3583	0.4945	0.1946	0.3254	0.3948	0.1557	0.2020	0.7980	0.8503	
6	7	104.0	428.0	89.0	83.0	187.0	517.0	0.5561	0.8279	0.1721	...	0.5389	0.5561	0.5474	0.2443	0.3840	0.3801	0.1562	0.1991	0.8009	0.8027	
7	8	71.0	491.0	26.0	116.0	187.0	517.0	0.3979	0.9497	0.0503	...	0.7320	0.3797	0.5000	0.2017	0.3294	0.3892	0.1560	0.2001	0.7999	0.8371	
8	9	103.0	472.0	45.0	84.0	187.0	517.0	0.5508	0.9130	0.0870	...	0.6959	0.5508	0.6149	0.1832	0.4638	0.4968	0.1581	0.1896	0.8104	0.8314	
9	10	98.0	453.0	64.0	89.0	187.0	517.0	0.5241	0.8762	0.1238	...	0.6049	0.5241	0.5616	0.2173	0.4003	0.4181	0.1541	0.2102	0.7898	0.8327	
11 rows × 23 columns																						

11 rows × 23 columns

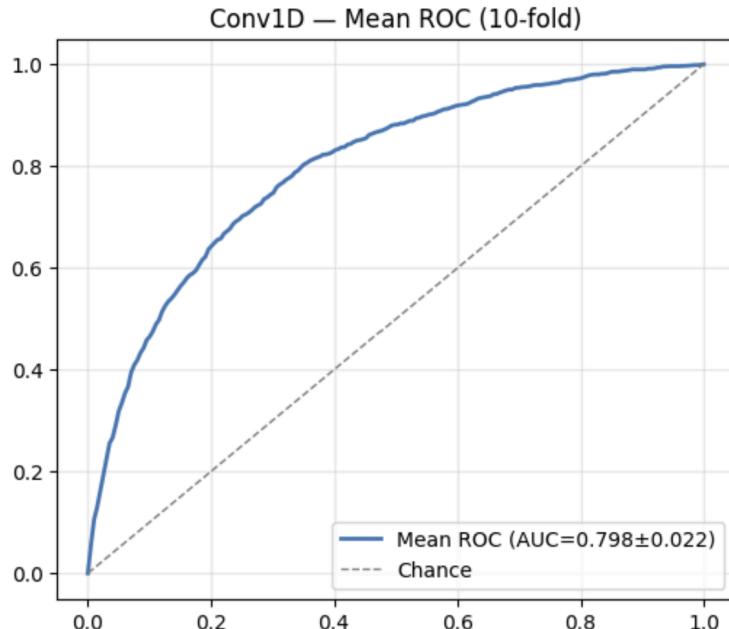
SVM (RBF) — Mean ROC (10-fold)



Per fold (Conv 1D) and ROC curve

== Conv1D (per fold + average) ==																								
	Fold	TP	TN	FP	FN	P	N	TPR	TNR	FPR	...	Precision	Recall	F1	ErrorRate	TSS	HSS	BS	BSS	BSS_alt	AUC			
0	1	131.0	417.0	101.0	56.0	187.0	518.0	0.7005	0.8050	0.1950	...	0.5647	0.7005	0.6253	0.2227	0.5056	0.4695	0.1655	0.1508	0.8492	0.8319			
1	2	124.0	407.0	111.0	63.0	187.0	518.0	0.6631	0.7857	0.2143	...	0.5277	0.6631	0.5877	0.2468	0.4488	0.4148	0.1673	0.1418	0.8582	0.8195			
2	3	128.0	405.0	113.0	59.0	187.0	518.0	0.6845	0.7819	0.2181	...	0.5311	0.6845	0.5981	0.2440	0.4663	0.4270	0.1865	0.0429	0.9571	0.8052			
3	4	107.0	429.0	89.0	79.0	186.0	518.0	0.5753	0.8282	0.1718	...	0.5459	0.5753	0.5602	0.2386	0.4035	0.3966	0.1710	0.1203	0.8797	0.8026			
4	5	121.0	443.0	74.0	66.0	187.0	517.0	0.6471	0.8569	0.1431	...	0.6205	0.6471	0.6335	0.1989	0.5039	0.4971	0.1792	0.0815	0.9185	0.8178			
5	6	112.0	423.0	94.0	75.0	187.0	517.0	0.5989	0.8182	0.1818	...	0.5437	0.5989	0.5700	0.2401	0.4171	0.4040	0.1916	0.0180	0.9820	0.7743			
6	7	120.0	376.0	141.0	67.0	187.0	517.0	0.6417	0.7273	0.2727	...	0.4598	0.6417	0.5357	0.2955	0.3690	0.3276	0.2001	-0.0259	1.0259	0.7596			
7	8	98.0	457.0	60.0	89.0	187.0	517.0	0.5241	0.8839	0.1161	...	0.6203	0.5241	0.5681	0.2116	0.4080	0.4293	0.1750	0.1031	0.8969	0.7761			
8	9	94.0	442.0	75.0	93.0	187.0	517.0	0.5027	0.8549	0.1451	...	0.5562	0.5027	0.5281	0.2386	0.3576	0.3689	0.1662	0.1477	0.8523	0.7890			
9	10	105.0	428.0	89.0	82.0	187.0	517.0	0.5615	0.8279	0.1721	...	0.5412	0.5615	0.5512	0.2429	0.3894	0.3848	0.1741	0.1077	0.8923	0.8032			
	Average	114.0	422.7	94.7	72.9	186.9	517.4	0.6099	0.8170	0.1830	...	0.5511	0.6099	0.5758	0.2380	0.4269	0.4120	0.1776	0.0888	0.9112	0.7979			

11 rows × 23 columns



Random Forest demonstrated the best overall balance between accuracy and recall, achieving an average accuracy of ~78.9%, F1 = 0.63, and AUC = 0.85. It maintained stable performance across folds with moderate variance and strong calibration (Brier Score ≈ 0.155).

SVM (RBF kernel) achieved the highest precision (0.60) and specificity (TNR ≈ 0.91), indicating it was more conservative and produced fewer false positives. However, it suffered from lower recall (0.47) and F1 (0.51), meaning it missed more true churners. Despite this, SVM's AUC (0.84) shows it still distinguished the two classes well.

Conv1D (Deep Learning) achieved balanced sensitivity and specificity (TPR ≈ 0.61, TNR ≈ 0.82) with an AUC of 0.80. It captured complex feature patterns but was slightly less stable across folds, possibly due to limited data size and training variability. Its Brier Skill Score (BSS ≈ 0.09) and F1 (0.57) indicate moderate predictive reliability.

7. Discussion

Algorithm Performance:

Among the three models, the Random Forest performed best overall, achieving the highest accuracy ($\approx 78.9\%$), F1-score (0.63), and AUC (0.85). Its ensemble structure of multiple decision trees helped reduce variance and overfitting while effectively handling both numerical and categorical variables. The model's ability to capture nonlinear feature interactions made it particularly well-suited for the Telco Churn dataset, which contains a mix of demographic and service-based attributes. The SVM (RBF kernel) showed strong precision and specificity but lower recall, suggesting it was more conservative and better at identifying non-churners but missed some true churn cases. The Conv1D neural network performed reasonably well but did not surpass Random Forest, likely due to the dataset's small size which limited the benefits of deep learning architectures that usually require larger training sets to generalize effectively.

Challenges and Solutions:

A major challenge was handling the mixed feature types (categorical and numerical) and missing values in TotalCharges. This was addressed by using One-Hot Encoding for categorical variables, StandardScaler for numerical ones, and imputing missing values with the median. The dataset was also slightly imbalanced ($\approx 26\%$ churners), so class weights were applied to all models and weighted loss for the Conv1D model to ensure fair learning. Another challenge was training time, especially with 10-fold cross-validation and hyperparameter tuning. This was mitigated by performing a single global RandomizedSearchCV (3-fold) before cross-validation to reuse optimized parameters efficiently.

Potential Improvements and Future Work:

Future work could focus on improving recall (churn detection) using SMOTE or ADASYN oversampling, threshold tuning. For deep learning, experimenting with LSTM or Bi-LSTM architectures, dropout optimization, or hyperparameter tuning (via grid or Bayesian search) could yield better results. Additionally, applying explainability techniques such as SHAP or LIME would help interpret feature importance and improve trust in model predictions.

8. Conclusion

This project successfully implemented and compared three binary classification algorithms: Random Forest, Conv1D Neural Network, and Support Vector Machine (SVM) to predict customer churn using the Telco Customer Churn dataset. The Random Forest consistently outperformed the other models across most metrics, achieving the highest accuracy ($\approx 79\%$), F1-score (0.63), and AUC (0.85), demonstrating strong generalization and robustness. The SVM achieved high precision and specificity, while Conv1D provided balanced performance but required more data to reach its full potential.

The findings highlight the impact of algorithm selection and evaluation methodology in predictive modeling. The use of 10-fold cross-validation ensured reliable and unbiased performance estimation, while comprehensive metrics such as AUC, F1, Brier Score, and Skill Scores provided a holistic view of model performance. Random Forest's ensemble nature offered the best trade-off between bias and variance for this mixed-type, moderately sized dataset.

Overall, the project demonstrates that ensemble models like Random Forest are well-suited for structured tabular data with both categorical and numerical features, whereas deep learning architectures may require larger datasets and more extensive tuning to outperform classical approaches. Future work can focus on feature engineering, data balancing, and model explainability to further improve churn prediction accuracy and interpretability.

9. References

Dataset Link: <https://www.kaggle.com/datasets/blastchar/telco-customer-churn>

Referred to Class Slides.

10. Github Link

My personal account: https://github.com/yashah9/shah_yash_finaltermproj

11. Appendix

Screenshot of Jupyter Notebook (training and evaluation):

```
(churn_env) PS C:\CS634-Final> jupyter notebook .\Yash_Shah_Final.ipynb
[I 2025-11-08 15:59:49.229 ServerApp] Extension package jupyter_lsp took 0.1422s to import
[I 2025-11-08 15:59:49.358 ServerApp] Extension package jupyter_server_terminals took 0.1282s to import
[I 2025-11-08 15:59:50.179 ServerApp] jupyter_lsp | extension was successfully linked.
[I 2025-11-08 15:59:50.186 ServerApp] jupyter_server_terminals | extension was successfully linked.
[I 2025-11-08 15:59:50.194 ServerApp] jupyterlab | extension was successfully linked.
[I 2025-11-08 15:59:50.201 ServerApp] notebook | extension was successfully linked.
[I 2025-11-08 15:59:50.634 ServerApp] notebook_shim | extension was successfully linked.
[I 2025-11-08 15:59:50.689 ServerApp] notebook_shim | extension was successfully loaded.
[I 2025-11-08 15:59:50.691 ServerApp] jupyter_lsp | extension was successfully loaded.
[I 2025-11-08 15:59:50.692 ServerApp] jupyter_server_terminals | extension was successfully loaded.
[I 2025-11-08 15:59:50.697 LabApp] JupyterLab extension loaded from C:\CS634-Final\churn_env\Lib\site-packages\jupyterlab
[I 2025-11-08 15:59:50.698 LabApp] JupyterLab application directory is C:\CS634-Final\churn_env\share\jupyter\lab
[I 2025-11-08 15:59:50.698 LabApp] Extension Manager is 'pypi'.
[I 2025-11-08 15:59:51.252 ServerApp] jupyterlab | extension was successfully loaded.
[I 2025-11-08 15:59:51.258 ServerApp] notebook | extension was successfully loaded.
[I 2025-11-08 15:59:51.260 ServerApp] Serving notebooks from local directory: C:\CS634-Final
[I 2025-11-08 15:59:51.261 ServerApp] Jupyter Server 2.17.0 is running at:
[I 2025-11-08 15:59:51.261 ServerApp] http://localhost:8888/tree?token=b5951aa2cd79ba13bed11c1d6670ac2c9cdb2cb9aae2f8c
[I 2025-11-08 15:59:51.261 ServerApp] http://127.0.0.1:8888/tree?token=b5951aa2cd79ba13bed11c1d6670ac2c9cdb2cb9aae2f8c
```

```

[GLOBAL] RF best params: {'clf__max_depth': 20, 'clf__max_features': 'log2', 'clf__min_samples_leaf': 9, 'clf__min_samples_split': 8, 'clf__n_estimators': 400}
[GLOBAL] SVM best params: {'clf__C': np.float64(21.368329072358772), 'clf__gamma': np.float64(0.0007068974950624604)}
Epoch 1/36
90/90      3s 12ms/step - loss: 0.6939 - val_loss: 0.6901 - learning_rate: 0.0010
Epoch 2/36
90/90      1s 12ms/step - loss: 0.6778 - val_loss: 0.6883 - learning_rate: 0.0010
Epoch 3/36
90/90      1s 13ms/step - loss: 0.6682 - val_loss: 0.6780 - learning_rate: 0.0010
Epoch 4/36
90/90      1s 13ms/step - loss: 0.6370 - val_loss: 0.5724 - learning_rate: 0.0010
Epoch 5/36
90/90      1s 8ms/step - loss: 0.6066 - val_loss: 0.6272 - learning_rate: 0.0010
Epoch 6/36
90/90      1s 7ms/step - loss: 0.5813 - val_loss: 0.5547 - learning_rate: 0.0010
Epoch 7/36
90/90      1s 7ms/step - loss: 0.5659 - val_loss: 0.6022 - learning_rate: 0.0010
Epoch 8/36
90/90      1s 8ms/step - loss: 0.5645 - val_loss: 0.5101 - learning_rate: 0.0010
Epoch 9/36
90/90      1s 8ms/step - loss: 0.5450 - val_loss: 0.5203 - learning_rate: 0.0010
Epoch 10/36
90/90      1s 8ms/step - loss: 0.5539 - val_loss: 0.5390 - learning_rate: 0.0010
Epoch 11/36
90/90      1s 8ms/step - loss: 0.5496 - val_loss: 0.5078 - learning_rate: 5.0000e-04
Epoch 12/36
90/90      1s 8ms/step - loss: 0.5520 - val_loss: 0.5273 - learning_rate: 5.0000e-04
Epoch 13/36
90/90      1s 7ms/step - loss: 0.5311 - val_loss: 0.5388 - learning_rate: 5.0000e-04
Epoch 14/36
90/90      1s 8ms/step - loss: 0.5200 - val_loss: 0.5483 - learning_rate: 2.5000e-04
Epoch 15/36
90/90      1s 8ms/step - loss: 0.5348 - val_loss: 0.5348 - learning_rate: 2.5000e-04
Epoch 1/36
90/90      3s 12ms/step - loss: 0.6881 - val_loss: 0.6872 - learning_rate: 0.0010
Epoch 2/36
90/90      1s 7ms/step - loss: 0.6852 - val_loss: 0.6855 - learning_rate: 0.0010
Epoch 3/36
90/90      1s 8ms/step - loss: 0.6910 - val_loss: 0.6384 - learning_rate: 0.0010
Epoch 4/36
90/90      1s 8ms/step - loss: 0.6656 - val_loss: 0.5793 - learning_rate: 0.0010
Epoch 5/36
90/90      1s 8ms/step - loss: 0.6237 - val_loss: 0.5274 - learning_rate: 0.0010
Epoch 6/36
90/90      1s 8ms/step - loss: 0.5893 - val_loss: 0.5387 - learning_rate: 0.0010
Epoch 7/36
90/90      1s 8ms/step - loss: 0.5654 - val_loss: 0.5148 - learning_rate: 0.0010
Epoch 8/36
90/90      1s 8ms/step - loss: 0.5652 - val_loss: 0.5499 - learning_rate: 0.0010
Epoch 9/36
90/90      1s 8ms/step - loss: 0.5601 - val_loss: 0.5519 - learning_rate: 0.0010
Epoch 10/36
90/90      1s 8ms/step - loss: 0.5617 - val_loss: 0.5072 - learning_rate: 5.0000e-04
Epoch 11/36
90/90      1s 8ms/step - loss: 0.5395 - val_loss: 0.5279 - learning_rate: 5.0000e-04
Epoch 12/36
90/90      1s 8ms/step - loss: 0.5453 - val_loss: 0.5077 - learning_rate: 5.0000e-04
Epoch 13/36
90/90      1s 8ms/step - loss: 0.5377 - val_loss: 0.5030 - learning_rate: 2.5000e-04
90/90      1s 8ms/step - loss: 0.5377 - val_loss: 0.5030 - learning_rate: 2.5000e-04

```

Evaluations are shown in results.

Screenshot of .py files (training and evaluation):

```

(churn_env) PS C:\CS634-Final> python ./churn_models.py
2025-11-10 22:23:32.270906: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0` .
2025-11-10 22:23:37.467084: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0` .
Shapes: (7043, 19) (7043,)
Positive rate (churn=1): 0.2653698707936959

==== Fold 1/10 ====
2025-11-10 22:23:47.739396: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
Epoch 1/36
90/90      2s 5ms/step - loss: 0.6924 - val_loss: 0.7031 - learning_rate: 0.0010
Epoch 2/36
90/90      0s 3ms/step - loss: 0.6888 - val_loss: 0.7050 - learning_rate: 0.0010
Epoch 3/36
90/90      0s 3ms/step - loss: 0.6733 - val_loss: 0.7268 - learning_rate: 0.0010
Epoch 4/36
90/90      0s 3ms/step - loss: 0.6490 - val_loss: 0.6564 - learning_rate: 5.0000e-04
Epoch 5/36
90/90      0s 3ms/step - loss: 0.6310 - val_loss: 0.5880 - learning_rate: 5.0000e-04
Epoch 6/36
90/90      0s 3ms/step - loss: 0.6155 - val_loss: 0.5999 - learning_rate: 5.0000e-04
Epoch 7/36
90/90      0s 3ms/step - loss: 0.6054 - val_loss: 0.5792 - learning_rate: 5.0000e-04
Epoch 8/36
90/90      0s 3ms/step - loss: 0.5993 - val_loss: 0.6183 - learning_rate: 5.0000e-04
Epoch 9/36
90/90      0s 4ms/step - loss: 0.5958 - val_loss: 0.5586 - learning_rate: 5.0000e-04
Epoch 10/36
90/90      0s 3ms/step - loss: 0.5886 - val_loss: 0.5683 - learning_rate: 5.0000e-04
Epoch 11/36
90/90      0s 3ms/step - loss: 0.5793 - val_loss: 0.5465 - learning_rate: 5.0000e-04
Epoch 12/36
90/90      0s 3ms/step - loss: 0.5769 - val_loss: 0.5832 - learning_rate: 5.0000e-04
90/90      0s 3ms/step - loss: 0.5769 - val_loss: 0.5832 - learning_rate: 5.0000e-04

```

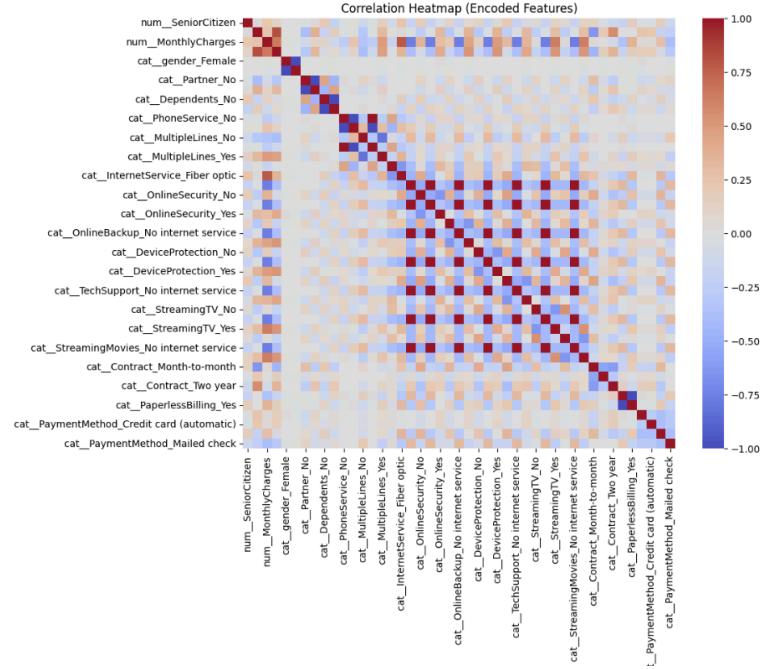
(churn_env) PS C:\CS634-Final> python .\churn_models.py											
	Fold	Accuracy	BalancedAcc	Precision	Recall	F1	TSS	HSS	BS	BSS	AUC
0	1	0.7773	0.7784	0.5573	0.7807	0.6503	0.5568	0.4936	0.1515	0.2227	0.8589
1	2	0.7702	0.7702	0.5475	0.7701	0.6400	0.5403	0.4782	0.1462	0.2499	0.8651
2	3	0.7589	0.7676	0.5307	0.7861	0.6336	0.5351	0.4638	0.1578	0.1903	0.8533
3	4	0.7571	0.7574	0.5281	0.7581	0.6225	0.5148	0.4518	0.1595	0.1797	0.8362
4	5	0.7869	0.7884	0.5714	0.7914	0.6637	0.5767	0.5136	0.1458	0.2528	0.8635
5	6	0.7699	0.7734	0.5468	0.7807	0.6432	0.5467	0.4810	0.1522	0.2199	0.8574
6	7	0.7315	0.7336	0.4964	0.7380	0.5935	0.4672	0.4044	0.1809	0.0724	0.7991
7	8	0.7955	0.7754	0.5931	0.7326	0.6555	0.5508	0.5123	0.1427	0.2687	0.8562
8	9	0.7500	0.7615	0.5194	0.7861	0.6255	0.5230	0.4494	0.1603	0.1782	0.8385
9	10	0.7599	0.7427	0.5366	0.7059	0.6097	0.4854	0.4410	0.1546	0.2073	0.8381
10	Average	0.7657	0.7648	0.5427	0.7630	0.6338	0.5297	0.4689	0.1551	0.2042	0.8466
==== SVM (10-fold) ====											
	Fold	Accuracy	BalancedAcc	Precision	Recall	F1	TSS	HSS	BS	BSS	AUC
0	1	0.8057	0.7516	0.6330	0.6364	0.6347	0.5032	0.5023	0.1389	0.2872	0.8566
1	2	0.8071	0.6842	0.7383	0.4225	0.5374	0.3684	0.4267	0.1559	0.1999	0.8465
2	3	0.8085	0.7603	0.6340	0.6578	0.6457	0.5207	0.5145	0.1470	0.2457	0.8559
3	4	0.8054	0.7316	0.6485	0.5753	0.6097	0.4633	0.4807	0.1583	0.1857	0.8291
4	5	0.7344	0.5000	0.0000	0.0000	0.0000	0.0000	0.0000	0.1527	0.2172	0.8649
5	6	0.8054	0.6627	0.7976	0.3583	0.4945	0.3254	0.3948	0.1557	0.2020	0.8503
6	7	0.7557	0.6920	0.5389	0.5561	0.5474	0.3840	0.3801	0.1562	0.1991	0.8027
7	8	0.7983	0.6647	0.7320	0.3797	0.5000	0.3294	0.3892	0.1560	0.2001	0.8371
8	9	0.8168	0.7319	0.6959	0.5508	0.6149	0.4638	0.4968	0.1581	0.1896	0.8314
9	10	0.7827	0.7001	0.6049	0.5241	0.5616	0.4003	0.4181	0.1541	0.2102	0.8327
10	Average	0.7920	0.6879	0.6023	0.4661	0.5146	0.3758	0.4003	0.1533	0.2137	0.8407
==== Conv1D (10-fold) ====											
	Fold	Accuracy	BalancedAcc	Precision	Recall	F1	TSS	HSS	BS	BSS	AUC
0	1	0.7844	0.7491	0.5806	0.6738	0.6238	0.4981	0.4738	0.1695	0.1305	0.8255
1	2	0.7872	0.7288	0.5979	0.6043	0.6011	0.4576	0.4560	0.1624	0.1666	0.8183
2	3	0.7447	0.7306	0.5137	0.7005	0.5928	0.4612	0.4132	0.1856	0.0475	0.8114
3	4	0.7628	0.6872	0.5537	0.5269	0.5399	0.3744	0.3803	0.1751	0.0991	0.7602
4	5	0.7997	0.7510	0.6173	0.6471	0.6319	0.5020	0.4944	0.1518	0.2220	0.8452
5	6	0.7713	0.7282	0.5613	0.6364	0.5965	0.4565	0.4378	0.1679	0.1394	0.8222
6	7	0.7003	0.6560	0.4487	0.5615	0.4988	0.3120	0.2888	0.2013	-0.0322	0.7366
7	8	0.7827	0.6848	0.6181	0.4759	0.5378	0.3696	0.3988	0.1603	0.1782	0.8034
8	9	0.7713	0.7180	0.5650	0.6043	0.5840	0.4360	0.4265	0.1675	0.1415	0.8063
9	10	0.7500	0.7206	0.5234	0.6578	0.5829	0.4411	0.4077	0.1721	0.1178	0.8137
10	Average	0.7654	0.7154	0.5580	0.6088	0.5789	0.4308	0.4177	0.1713	0.1210	0.8043
==== Summary (Averages Across Folds) ====											
	Model	Accuracy	BalancedAcc	Precision	Recall	F1	TSS	HSS	BS	BSS	AUC
0	RandomForest	0.7657	0.7648	0.5427	0.7630	0.6338	0.5297	0.4689	0.1551	0.2042	0.8466
1	SVM_RBF	0.7920	0.6879	0.6023	0.4661	0.5146	0.3758	0.4003	0.1533	0.2137	0.8407
2	Conv1D	0.7654	0.7154	0.5580	0.6088	0.5789	0.4308	0.4177	0.1713	0.1210	0.8043

Additional Plots:

I have only built plots in the Jupyter Notebook and not in .py files because I was having trouble displaying them.

ROC curves are displayed in the Results section.

Encoded feature matrix shape: (7043, 45)
Example columns: ['num_SeniorCitizen', 'num_tenure', 'num_MonthlyCharges', 'num_TotalCharges', 'cat_gender_Female', 'cat_gender_Male', 'cat_Partner_No', 'cat_Partner_Yes', 'cat_Dependents_No', 'cat_Dependents_Yes']



Top 10 Important Features for Churn Prediction

